

Mid2 Report

LI XU

April 18, 2018

1 Modifications

Similar to the cifar10_cnn example. I use image_data_generator function to enable the image random shifting and rotating when training. Also I set the epochs to 5. The model fit code is below:

```
datagen = image_data_generator(rotation_range=8,
width_shift_range=0.08,
shear_range=0.3,
height_shift_range=0.08,
zoom_range=0.08)
datagen %>% fit_image_data_generator(x_train)
model %>% fit_generator(
  flow_images_from_data(x_train, y_train, datagen, batch_size = batch_
    size),
  steps_per_epoch = as.integer(60000/64),
  epochs = 5,
  validation_data = list(x_test, y_test),
  validation_steps=10000/64
)
```

2 Results

```
Epoch 1/5
937/937 [=====] - 19s 20ms/step - loss:
0.0599 - acc: 0.9832 - val_loss: 0.0267 - val_acc: 0.9927
Epoch 2/5
937/937 [=====] - 19s 20ms/step - loss:
0.0609 - acc: 0.9830 - val_loss: 0.0185 - val_acc: 0.9948
Epoch 3/5
937/937 [=====] - 19s 20ms/step - loss:
0.0595 - acc: 0.9832 - val_loss: 0.0208 - val_acc: 0.9934
Epoch 4/5
937/937 [=====] - 19s 20ms/step - loss:
0.0589 - acc: 0.9836 - val_loss: 0.0186 - val_acc: 0.9948
Epoch 5/5
```

```

937/937 [=====] - 19s 20ms/step - loss:
0.0600 - acc: 0.9835 - val_loss: 0.0188 - val_acc: 0.9947
> scores <- model %>% evaluate(
+   x_test, y_test, verbose = 0
+ )
> cat('Test_loss:', scores[[1]], '\n')
Test loss: 0.0187528
> cat('Test_accuracy:', scores[[2]], '\n')
Test accuracy: 0.9947

```

The figure during training is in Figure (1). The testing accuracy is 99.47%




Figure 1: Results plot.

A R code

```
library(keras)

# Data Preparation 

---



batch_size <- 128
num_classes <- 10
epochs <- 12

# Input image dimensions
img_rows <- 28
img_cols <- 28

# The data, shuffled and split between train and test sets
mnist <- dataset_mnist()
x_train <- mnist$train$x
y_train <- mnist$train$y
x_test <- mnist$test$x
y_test <- mnist$test$y

# Redefine dimension of train/test inputs
x_train <- array_reshape(x_train, c(nrow(x_train), img_rows, img_cols, 1))
x_test <- array_reshape(x_test, c(nrow(x_test), img_rows, img_cols, 1))
input_shape <- c(img_rows, img_cols, 1)

# Transform RGB values into [0,1] range
x_train <- x_train / 255
x_test <- x_test / 255

cat('x_train_shape:', dim(x_train), '\n')
cat(nrow(x_train), 'train_samples\n')
cat(nrow(x_test), 'test_samples\n')

# Convert class vectors to binary class matrices
y_train <- to_categorical(y_train, num_classes)
y_test <- to_categorical(y_test, num_classes)

# Define Model 

---



# Define model
model <- keras_model_sequential() %>%
layer_conv_2d(filters = 32, kernel_size = c(3,3), activation = 'relu',
input_shape = input_shape) %>%
layer_conv_2d(filters = 64, kernel_size = c(3,3), activation = 'relu') %>%
layer_max_pooling_2d(pool_size = c(2, 2)) %>%
layer_dropout(rate = 0.25) %>%
layer_flatten() %>%
layer_dense(units = 128, activation = 'relu') %>%
```

```

layer_dropout(rate = 0.5) %>%
layer_dense(units = num_classes, activation = 'softmax')
model %>% compile(
loss = loss_categorical_crossentropy,
optimizer = optimizer_adadelta(),
metrics = c('accuracy')
)

datagen = image_data_generator(rotation_range=8,
width_shift_range=0.08,
shear_range=0.3,
height_shift_range=0.08,
zoom_range=0.08)
datagen %>% fit_image_data_generator(x_train)
model %>% fit_generator(
flow_images_from_data(x_train, y_train, datagen, batch_size = batch_size),
steps_per_epoch = as.integer(60000/64),
epochs = 5,
validation_data = list(x_test, y_test),
validation_steps=10000/64
)

scores <- model %>% evaluate(
x_test, y_test, verbose = 0
)

# Output metrics
cat('Test_loss:', scores[[1]], '\n')
cat('Test_accuracy:', scores[[2]], '\n')

```